

# Chat with Multiple PDFs using Ollama

Date: 15th September 2025

---

## ☐ Task Description

An **AI-powered document assistant** that allows users to chat with multiple PDFs. It combines **RAG (Retrieval-Augmented Generation)** with **Ollama's Gemma3:4b LLM** to make document understanding more interactive and user-friendly.

---

## ⚡ Key Features Implemented

- Uploading multiple PDFs
  - Extracting and chunking text
  - Creating embeddings using **Sentence Transformers (all-MiniLM-L6-v2)**
  - Storing/retrieving data with **ChromaDB**
  - Querying the knowledge base with **semantic search**
  - Generating contextual answers via **Ollama (Gemma3:4b)**
  - Building a **Streamlit UI** for document Q&A
- 

## ☐ Activities (Timeline)

- **10:30 – 11:30** → R&D and installation of Gemma model.
  - **11:30 – 1:00** → Extracted text from PDFs and generated answers using Gemma3:4b.
  - **1:30 – 2:30** → Lunch break.
  - **2:30 – 3:30** → Learned FAISS indexing and Qdrant vector DB concepts.
  - **3:30 – 5:30** → Developed the RAG model integrated with Ollama; implemented Streamlit UI for multi-document chat. Encountered bugs that require fine-tuning.
  - **5:30 – 6:30** → Discussion with Rajasekar Anna about project scope and AI trends.
  - **6:45** → Left for the day.
- 

## ☐ Skills & Learnings

## From Today's Work:

- Handling structured/unstructured text extraction from PDFs with **PyPDF2**.
- Understanding embeddings and semantic similarity using **Sentence Transformers**.
- Exploring **ChromaDB** collections, embedding storage, and query operations.
- Implementing a complete **RAG pipeline with LangChain + Ollama**.
- Building a user-friendly **Streamlit interface** with sidebar file upload & chat input.
- API integration with **Ollama's LLM** for contextual responses.
- Troubleshooting **Python version issues** and database refresh logic.

## Key Takeaways:

- RAG systems are only as good as the **quality of embeddings and retrieval**.
  - Choosing the right **vector DB (FAISS, Qdrant, Chroma)** impacts scalability and efficiency.
  - **UI matters**: an interactive frontend makes backend functionality much more accessible.
- 

## □ Challenges & Solutions


### Obstacles Faced:

- Python version conflicts and dependency mismatches.
- Initial implementation only worked for **single documents**.
- Errors when scaling to **multi-document setup**.

### Strategies Adopted:

- Did **step-by-step R&D** with single-document pipeline before moving to DB-based multi-document setup.
  - Switched to **ChromaDB** for embedding storage.
  - Implemented **checks for empty/duplicate files** and state refresh in Streamlit.
- 

## Output :



# □ Self-Reflection

## Observations:

Today gave me a **solid hands-on understanding** of how RAG pipelines work end-to-end. While theory (embeddings, retrieval) is clear, the real challenge was implementation—dealing with **bugs, DB handling, and Python compatibility**.

The discussion with **Rajasekar Anna** gave useful clarity on scaling this project.

## Future Action:

- Implement **FAISS** and **Qdrant** hands-on for comparison.
- Debug **multi-document retrieval accuracy** and Ollama integration.
- Enhance **Streamlit UI** with expandable retrieved context.
- Explore **prompt engineering tweaks** for more precise answers.